

A Stochastic Conjugate Algorithms with the Minimized Variance Reduction

Hongbo Du

Instructed by Caixia Kou

School of Science
Beijing University of Posts and Telecommunications

December 1, 2023

Outline

- ① Introduction
- ② The Minimal Variance Stochastic Gradient Estimate
- ③ Algorithm and Convergence analysis
- ④ Numerical Experiments
- ⑤ Conclusion

- ① Introduction
- ② The Minimal Variance Stochastic Gradient Estimate
- ③ Algorithm and Convergence analysis
- ④ Numerical Experiments
- ⑤ Conclusion

Introduction

The research problem

$$\min_{\omega \in R^d} f(\omega) = \frac{1}{n} \sum_{i=1}^n f_i(\omega)$$

Linear regression (Ridge regression)

$$\min_{\omega \in R^d} f(\omega) = \frac{1}{n} \sum_{i=1}^n (x_i^T \omega - y_i)^2 + \frac{\lambda}{2} \|\omega\|^2 = \frac{1}{n} \sum_{i=1}^n f_i(\omega)$$

$$f_i(\omega) = (x_i^T \omega - y_i)^2 + \frac{\lambda}{2} \|\omega\|^2$$

Logistic regression

$$\min_{\omega \in R^d} f(\omega) = \frac{1}{n} \sum_{i=1}^n \log(1 + \exp(-y_i x_i^T \omega)) + \frac{\lambda}{2} \|\omega\|^2$$

$$f_i(\omega) = \log(1 + \exp(-y_i x_i^T \omega)) + \frac{\lambda}{2} \|\omega\|^2$$

Introduction

GD(Gradient Decent)

$$\omega_{k+1} = \omega_k - \alpha_k \nabla f(\omega_k)$$

$$\nabla f(\omega_k) = \frac{1}{n} \sum_{i=1}^n \nabla f_i(\omega_k)$$

SGD(Robbin, Monro, 1951)

$$\omega_{k+1} = \omega_k - \alpha_k \nabla f_i(\omega_k)$$

Mini-batch SGD (Shalev-Shwartz et al., 2007)

Fixed stepsize:

$$\omega_{k+1} = \omega_k - \alpha \nabla f_{S_k}(\omega_k)$$

Decreased stepsize:

$$\omega_{k+1} = \omega_k - \alpha_k \nabla f_{S_k}(\omega_k)$$

where $\nabla f_{S_k}(\omega_k) = \frac{1}{|S|} \sum_{i \in S_k} \nabla f_i(\omega_k)$

The Stochastic Gradient Algorithm with Variance Reduction

SVRG(R. Johnson and T. Zhang, 2013)

$$g_k = \nabla f_j(\omega^k) - \nabla f_j(\psi_j^k) + \frac{1}{n} \sum_{i=1}^n \nabla f_i(\psi_i^k)$$

$$g_k = \nabla f_{S_k}(\omega^k) - \nabla f_{S_k}(\psi_j^k) + \frac{1}{n} \sum_{i=1}^n \nabla f_i(\psi_i^k)$$

SAGA(A. Defazio, F. Bach, and S. Lacoste-Julien, 2014)

$$g_k = \nabla f_{S_k}(\omega^k) - \nabla f_{S_k}(\phi_j^k) + \frac{1}{n} \sum_{i=1}^n \nabla f_i(\phi_i^k)$$

SAG (M. Schmidt, N. Le Roux, and F. Bach, 2017)

$$g_k = \frac{\nabla f_{S_k}(\omega^k) - \nabla f_{S_k}(\phi_j^k)}{n} + \frac{1}{n} \sum_{i=1}^n \nabla f_i(\phi_i^k)$$

- 1 Introduction
- 2 The Minimal Variance Stochastic Gradient Estimate
- 3 Algorithm and Convergence analysis
- 4 Numerical Experiments
- 5 Conclusion

Motivation

Let

$$\bar{X} = \nabla f_{S_k}(\omega^k), \bar{Y} = \nabla f_{S_k}(\phi_i^k), \text{ or } \bar{Y} = \nabla f_{S_k}(\psi_i^k)$$

(SAGA/SVRG)

$$g_k = \bar{X} - \bar{Y} + E(\bar{Y})$$

$$E(g_k) = E(\bar{X} - \bar{Y} + E(\bar{Y})) = E(\bar{X})$$

$$\text{Var}(g_k) = \text{Var}(\bar{X} - \bar{Y} + E(\bar{Y})) = \text{Var}(\bar{X} - \bar{Y})$$

(SAG)

$$g_k = \frac{\bar{X} - \bar{Y}}{n} + E(\bar{Y})$$

$$E(g_k) = E\left(\frac{\bar{X} - \bar{Y}}{n} + E(\bar{Y})\right) = \frac{1}{n}E(\bar{X}) + \left(1 - \frac{1}{n}\right)E(\bar{Y})$$

$$\text{Var}(g_k) = \text{Var}\left(\frac{\bar{X} - \bar{Y}}{n} + E(\bar{Y})\right) = \frac{1}{n^2} \text{Var}(\bar{X} - \bar{Y})$$

Gradient Estimate with Unbiasedness and Minimal Variance

Estimator:

$$\theta_\gamma = \bar{X} - \gamma(\bar{Y} - E(\bar{Y}))$$

Expected Value:

$$E(\theta_\gamma) = E(\bar{X})$$

Variance:

$$\begin{aligned} \text{Var}(\theta_\gamma) &= \text{Var}(\bar{X} - \gamma(\bar{Y} - E(\bar{Y}))) \\ &= \text{Var}(\bar{X}) + \gamma^2 \text{Var}(\bar{Y}) - 2\gamma \text{Cov}(\bar{X}, \bar{Y}) \end{aligned} \quad (1)$$

Best Param:

$$\gamma^* = \frac{\text{Cov}(\bar{X}, \bar{Y})}{\text{Var}(\bar{Y})} \approx \frac{s_{XY}}{s_Y^2} \quad (2)$$

$$\begin{aligned} s_{XY} &= \frac{1}{|S| - 1} \sum_{j \in S_k} (X_j - \bar{X})(Y_j - \bar{Y}) \\ s_Y^2 &= \frac{1}{|S| - 1} \sum (Y_j - \bar{Y})^2 \end{aligned} \quad (3)$$

Estimation of γ^*

$$\begin{aligned} \text{Cov}(\bar{X}, \bar{Y}) &= \text{Cov}\left(\frac{1}{|S|} \sum_{j \in S_k} X_j, \frac{1}{|S|} \sum_{j \in S_k} Y_j\right) = \frac{1}{|S|^2} \text{Cov}\left(\sum_{j \in S_k} X_j, \sum_{j \in S_k} Y_j\right) \\ &= \frac{1}{|S|^2} \sum_{j \in S_k} \text{Cov}(X_j, Y_j) = \frac{1}{|S|} \text{Cov}(X, Y) \approx \frac{1}{|S|} s_{XY} \end{aligned} \quad (4)$$

$$\begin{aligned} \text{Var}(\bar{Y}) &= \text{Var}\left(\frac{1}{|S|} \sum_{j \in S_k} Y_j\right) = \frac{1}{|S|^2} \text{Var}\left(\sum_{j \in S_k} Y_j\right) \\ &= \frac{1}{|S|^2} \sum_{j \in S_k} \text{Var}(Y_j) = \frac{1}{|S|} \text{Var}(Y) \approx \frac{1}{|S|} s_Y^2 \end{aligned} \quad (5)$$

$$\begin{aligned} s_{XY} &= \frac{1}{|S| - 1} \sum_{j \in S_k} (X_j - \bar{X})(Y_j - \bar{Y}) \\ s_Y^2 &= \frac{1}{|S| - 1} \sum_{j \in S_k} (Y_j - \bar{Y})^2 \end{aligned} \quad (6)$$

The Stochastic Conjugate Gradient Algorithm with Variance Reduction

CGVR(Ran Xin, Usman A. Khan, and Soummya Kar, 2020)

$$g_k = \nabla f_{S_k}(\omega^k) - \nabla f_{S_k}(\psi_j^k) + \frac{1}{n} \sum_{i=1}^n \nabla f_i(\psi_i^k)$$

$$d_k = -g_k + \beta_k d_{k-1}$$

SCGA(Caixia Kou and Han Yang, 2022)

$$g_k = \nabla f_{S_k}(\omega^k) - \nabla f_{S_k}(\phi_j^k) + \frac{1}{n} \sum_{i=1}^n \nabla f_i(\phi_i^k)$$

$$d_k = -g_k + \beta_k d_{k-1}$$

- ① Introduction
- ② The Minimal Variance Stochastic Gradient Estimate
- ③ Algorithm and Convergence analysis
- ④ Numerical Experiments
- ⑤ Conclusion

Algorithm 1: SCGA with the minimal variance stochastic gradient estimate

```

1 Initialization: .....
2 Iteration:
3 for  $k = 1, 2 \dots$  do
4     .....
5     for  $j : S_k$  do
6         Compute  $\nabla f_j(\omega_k)$  and store it into matrix  $\nabla f_{[S_k]}(\omega_k)$ 
7         Select  $\nabla f_j(\omega_{[k-1]})$  in  $\nabla f(\omega_{[k-1]})$  and store it into matrix
             $\nabla f_{[S_k]}(\omega_{[k-1]})$ 
8     for  $r = 1, 2 \dots d$  do
9         Compute the sample covariance of  $\nabla f_{[S_k]}^{(r)}(\omega_k)$  and
             $\nabla f_{[S_k]}^{(r)}(\omega_{[k-1]})$ , the sample variance of  $\nabla f_{[S_k]}^{(r)}(\omega_{[k-1]})$  using
            (6)
10        Compute  $\gamma^{*(r)}$  using (2)
11    Compute

```

Algorithm 2: CGVR with the minimal variance stochastic gradient estimate

```

1 Initialization: Given  $x_0 \in R^d$ , compute  $h_0 = \frac{1}{n} \sum_{i=1}^n \nabla f_i(x_0)$ :
2 Iteration: for  $l = 1, 2 \dots T$  do
3     .....
4     for  $k = 1, 2 \dots m$  do
5         .....
6         for  $j: S_k$  do
7             Compute  $\nabla f_j(\omega_k), \nabla f_j(\omega_0)$ 
8             Store  $\nabla f_{[S_k]}(\omega_k) \leftarrow \nabla f_j(\omega_k), \nabla f_{[S_k]}(\omega_0) \leftarrow \nabla f_j(\omega_0)$ 
9         for  $r = 1, 2 \dots d$  do
10             Compute the sample covariance  $\nabla f_{[S_k]}^{(r)}(\omega_k)$  and
11                  $\nabla f_{[S_k]}^{(r)}(\omega_0)$ , the sample variance of  $\nabla f_{[S_k]}^{(r)}(\omega_0)$  using (6)
12             Compute  $\gamma^{*(r)}$  using (2)
13         Compute  $\nabla f_{S_k}(\omega_k), \nabla f_{S_k}(\omega_0)$ 
14         Compute  $g_l = \nabla f_{S_k}(\omega_k) - \gamma^{*}(\nabla f_{S_k}(\omega_k) - \nabla f_{S_k}(\omega_0))$ 

```

Assumptions

Assumption 1 (μ -strong convexity and L -smoothness)

$f_i, 1 \leq i \leq n$ is strongly convex and has Lipschitz continuous gradients, i.e.,

$$\mu I \prec \nabla^2 f_i(w) \prec LI \quad (7)$$

For $\omega \in R^d$, μ is strong convexity constant and L is Lipschitz constant.

Assumption 2 (low and upper bounds of step size) every step size α_k in Alg1 and Alg2 algorithm satisfies $\alpha_1 \leq \alpha_k \leq \alpha_2$

Assumption 3 (upper bound of scalar β_k) There exists constants β such that

$$\beta_k \leq \frac{\|g_k\|^2}{\|g_{k-1}\|^2} \leq \beta \quad (8)$$

Related Lemma

Lemma

Under Assumption1, we have

$$2\mu(f(\omega) - f(\omega^*)) \leq \|\nabla f(w)\|^2 \leq 2L(f(\omega) - f(\omega^*)) \quad (9)$$

Where $\omega \in R^d$, ω^ is the unique minimizer*

Lemma

Consider that Alg1 and Alg2 (CG) algorithm, where step size α_k satisfies strong Wolfe condition with $0 < \sigma_2 < \frac{1}{2}$ and β_k satisfies $|\beta_k| \leq \beta_k^{FR}$, then it generates descent directions d_k satisfying

$$-\frac{1}{1 - \sigma_2} \leq \frac{\langle g_k, d_k \rangle}{\|g_k\|^2} \leq \frac{2\sigma_2 - 1}{1 - \sigma_2} \quad (10)$$

Convergence of Alg1

Theorem

Let Assumption 1,2,3 hold. If the bound of the step-size in Alg1 satisfies:

$$0 < \alpha_1 < \frac{1 - \beta}{2L\sigma_1} \quad (11)$$

Then we have: $\forall k > 0$

$$E(f(\omega_k)) - f(\omega^*) \leq C\xi^k(E(f(\omega_0)) - f(\omega^*)) \quad (12)$$

where $\xi = \frac{(1-\sigma_1)(1-\beta)+2L\alpha\sigma_1\sigma_2(1-\beta^m)}{2\mu\sigma_1m(1-\sigma_2)(1-\beta)} < 1$, ω^ is the unique minimizer of f*

Convergence of Alg2

Theorem

Let Assumption 1,2,3 hold. Let

$$m > \frac{(1 - \sigma_1) + 2L\alpha_2\sigma_1\sigma_2\beta}{2\mu\sigma_1\alpha_1(1 - \sigma_2)}$$

Then we have: $\forall l > 0$

$$E(f(x_l)) - f(\omega^*) \leq \xi^l (E(f(x_0)) - f(\omega^*)) \quad (13)$$

where

$$\xi = \frac{(1 - \sigma_1)(1 - \beta) + 2L\alpha\sigma_1\sigma_2(1 - \beta^m)}{2\mu\sigma_1m(1 - \sigma_2)(1 - \beta)} < 1$$

and ω^ is the unique minimizer of f*

- 1 Introduction
- 2 The Minimal Variance Stochastic Gradient Estimate
- 3 Algorithm and Convergence analysis
- 4 Numerical Experiments**
- 5 Conclusion

Datasets

Table: Summary of data sets used in numerical experiments

dataset	d	n	type
A9a	123	32561	binary classification
ljcnn1	22	49990	binary classification
Protein	74	145751	binary classification
Quantum	78	50000	binary classification
W8a	300	49749	binary classification
Covtype	54	581012	binary classification
YearPredictionMSD	90	463715	regression
Pyrim	27	74	regression
Bodyfat	24	252	regression
Triazines	60	180	regression
Eunite2001	16	336	regression

Test Function

The ridge regression model are presented as follows:

$$\min_{\omega} \frac{1}{n} \sum_{i=1}^n (y_i - x_i \omega)^2 + \lambda \|\omega\|^2 \quad (14)$$

where $x_i \in R^d$ is denoted the feature vector of the i-th data sample, $y_i \in R$ is denoted the actual value of the i-th data sample, and λ is the regularization parameter.

Numerical Results

Figure: Variance comparison of stochastic gradient estimates: $g_{S_l}^k(\gamma = \gamma^*)$ and $g_{S_l}^k(\gamma = 1)$

$$g_{S_l}^k(\gamma) = \nabla f_{S_l}(w_{101}) - \gamma(\nabla f_{S_l}(w_k) - \frac{1}{n} \sum_{i=1}^n \nabla f_i(w_k)) \quad l = 1, \dots, 100 \quad (15)$$

$$Var(g_{S_l}^k(\gamma)) = \frac{1}{100} \sum_{l=1}^{100} (g_{S_l}^k(\gamma) - \frac{1}{100} \sum_{l=1}^{100} g_{S_l}^k(\gamma))^2, \quad k = 0, 1, 2, \dots, 100 \quad (16)$$

Numerical Results

Figure: performance profiles of SCGA,CGVR,Alg1,Alg2 on the six data sets of binary classification (x-axis is times of iteration, y-axis is loss value in terms of \log_{10})

Numerical Results

Figure: performance profiles of SCGA,CGVR,Alg1,Alg2 on the six data sets of regression (x-axis is times of iteration, y-axis is loss value in terms of \log_{10})

- ① Introduction
- ② The Minimal Variance Stochastic Gradient Estimate
- ③ Algorithm and Convergence analysis
- ④ Numerical Experiments
- ⑤ Conclusion

Conclusion

- ① Propose the minimal variance stochastic gradient estimate
- ② Propose two improved algorithms of SCGA and CGVR: Alg1 and Alg2
- ③ Prove the linear convergence rate of the new algorithms under strong convexity and smoothness
- ④ From a series of experiments, compared with SCGA and CGVR, Alg1 and Alg2 is competitive algorithms.

Thanks for your attention!